Semiprime Bottlenecks in Arithmetic Graphs: Betweenness Centrality from Prime Highways

Tobias Canavesi

Independent Researcher

tobias.canavesi@crosschq.com

Abstract

We introduce a family of deterministic networks G_N on the integers $\{1, \ldots, N\}$, where nodes are linked by unit increments and by prime-stride shortcuts. Despite their simple construction, these *arithmetic graphs* exhibit a striking structural paradox: the nodes with the highest degree—primorials like $2310 = 2 \cdot 3 \cdot 5 \cdot 7 \cdot 11$—are *not* the most central. Instead, betweenness centrality is dominated by semiprimes of the form $2p$ and $3p$, where p is a large prime. We prove an exact degree formula showing $\deg(v) = 2 + 2\,\omega(v)$ for interior nodes (ω = number of distinct prime divisors), prove $\mathrm{diam}(G_N) = O(\sqrt{N})$ with empirical evidence supporting $O(\log N)$, and give a routing-theoretic explanation for why minimal-multiplier semiprimes maximize betweenness. Computations up to $N = 10{,}000$ confirm that over 96% of the top-50 betweenness nodes are semiprimes $2p$ or $3p$, while the degree–betweenness correlation decays below 0.03.

Keywords: arithmetic graph, betweenness centrality, prime numbers, semiprime, small-world network, routing

1 Introduction

The interplay between number theory and network topology is a source of surprising structural phenomena. In this paper, we study a deterministic graph on the integers where edges arise from two natural operations: sequential adjacency ($n \to n+1$) and prime-stride multiplication ($kp \to (k+1)p$ for each prime p). The resulting graph has a backbone chain augmented by "highways" at every prime scale.

A naive prediction would place the structurally most important nodes at the integers with the most prime factors—the highly composite numbers—since these sit on the most highways and have the highest degree. We show that this prediction is wrong. Betweenness centrality, the standard measure of a node's role as a routing bottleneck [1], is instead maximized at semiprimes $2p$ and $3p$, where p is a large prime.

The explanation is a trade-off between two competing costs. A large prime p creates a sparse highway with only $\lfloor N/p \rfloor$ stops, making each stop a critical bottleneck. Among these stops, the one at $2p$ is the most accessible because every node is at most one step from a multiple of 2. This *scarcity–accessibility trade-off* is the central mechanism we formalize.

2 The arithmetic graph

Definition 2.1 (Arithmetic graph)**.** For $N \geq 2$, the *arithmetic graph* $G_N = (V, E)$ has vertex set $V = \{1, 2, \ldots, N\}$ and edge set

$$E \;=\; \underbrace{\{(n, n+1) : 1 \leq n < N\}}_{\text{additive edges}} \cup \underbrace{\{(kp,\, (k+1)p) : p \text{ prime},\; k \geq 1,\; (k+1)p \leq N\}}_{\text{multiplicative edges}}. \tag{1}$$

The additive edges form a path graph (the "backbone"). For each prime p, the multiplicative edges link consecutive multiples of p, forming a *p-highway* with $\lfloor N/p \rfloor$ stops and step size p. The 2-highway connects every other integer, the 3-highway every third, and so on. A node v sits on the p-highway if and only if $p \mid v$.

Proposition 2.2 (Edge count)**.** *The number of edges in G_N is*

$$|E(G_N)| \;=\; (N-1) \;+\; \sum_{\substack{p \leq N \\ p \text{ prime}}} \Big(\Big\lfloor \frac{N}{p} \Big\rfloor - 1 \Big). \tag{2}$$

The average degree satisfies $\overline{\deg} \sim 2 \log\log N$ as $N \to \infty$.

Proof. The $N-1$ additive edges form a path. For each prime $p \leq N$, the p-highway has $\lfloor N/p \rfloor$ nodes and $\lfloor N/p \rfloor - 1$ edges. These edge sets are pairwise disjoint: a multiplicative edge $(kp, (k+1)p)$ has step size p, and distinct primes produce distinct step sizes. Moreover, no multiplicative edge is also additive, since $p \geq 2$ implies $(k+1)p - kp = p \geq 2 > 1$. Summing gives (2).

For the average degree, $\overline{\deg} = 2|E|/N$ and

$$\frac{|E|}{N} \;=\; 1 + \frac{1}{N}\sum_{p\le N}\left\lfloor \frac{N}{p}\right\rfloor - \frac{\pi(N)}{N} \;\sim\; \sum_{p\le N}\frac{1}{p} \;\sim\; \log\log N$$

by the Mertens estimate [6] $\sum_{p\le x} 1/p = \log\log x + M + o(1)$, where $M \approx 0.2615$ is the Meissel–Mertens constant. □

Remark 2.3. The graph is thus ultra-sparse (density $\to 0$) yet, as we show below, small-world in the sense of Watts and Strogatz [2] (diameter $O(\sqrt{N})$ proven, $O(\log N)$ empirically). The average degree diverges, but only at the doubly-logarithmic rate $\sim 2\log\log N$. This creates a regime where local connectivity grows very slowly while global shortcuts (from large-prime highways) provide efficient long-range navigation.

3 Degree structure

Proposition 3.1 (Degree formula)**.** *For any $v \in V$, write $v = \prod_{i=1}^{r} p_i^{a_i}$ with $p_1 < \cdots < p_r$ the distinct prime factors of v. Then*

$$\deg(v) \;=\; \mathbf{1}_{v>1} + \mathbf{1}_{v<N} + \sum_{p\mid v}\big(\mathbf{1}_{v-p\ge 1} + \mathbf{1}_{v+p\le N}\big). \tag{3}$$

For interior nodes (i.e., $v > p_r$ and $v + p_r \le N$), this simplifies to

$$\deg(v) = 2 + 2\,\omega(v), \tag{4}$$

where $\omega(v)$ denotes the number of distinct prime divisors of v (see [4], §22.1).

Proof. The additive edges contribute one neighbor on each side when v is interior ($1 < v < N$). For each prime $p \mid v$, node $v = kp$ sits on the p-highway between $(k-1)p$ and $(k+1)p$. The edge to $(k-1)p = v - p$ exists iff $v - p \geq 1$ (i.e., $k \geq 2$), and the edge to $(k+1)p = v + p$ exists iff $v + p \leq N$. No two of these neighbors coincide: the additive neighbors $v \pm 1$ are coprime to any prime $p \geq 2$ dividing v (since $\gcd(v, v \pm 1) = 1$), and distinct prime factors yield distinct highway neighbors ($v - p_i \neq v - p_j$ for $p_i \neq p_j$). □

Remark 3.2. The formula creates a clean hierarchy: primes ($\omega = 1$) at boundary positions have deg $= 3$, semiprimes kp ($\omega = 2$) have deg $= 6$, and primorials like $2310 = 2 \cdot 3 \cdot 5 \cdot 7 \cdot 11$ ($\omega = 5$) have deg $= 12$. If degree determined importance, the primorials would be the hubs. The data show otherwise.

Table 1: Degree formula verification.

v	Factorization	$\omega(v)$	$\deg(v)$	Predicted
106	$2 \cdot 53$	2	6	$2 + 2 \cdot 2 = 6$
173	prime	1	3	boundary: $2 + 1 = 3$
210	$2 \cdot 3 \cdot 5 \cdot 7$	4	10	$2 + 2 \cdot 4 = 10$
331	prime	1	3	boundary: $2 + 1 = 3$
2310	$2 \cdot 3 \cdot 5 \cdot 7 \cdot 11$	5	12	$2 + 2 \cdot 5 = 12$

4 Small-world property

Proposition 4.1. $\mathrm{diam}(G_N) = O(\sqrt{N})$.

Proof. We show that any node s can reach any node t in $O(\sqrt{N})$ steps via an iterated strategy using highways of decreasing stride.

Let $d = |s - t|$ be the remaining distance. Choose a prime p with $\lfloor \sqrt{d} \rfloor \leq p \leq 2\sqrt{d}$ (such a prime exists by Bertrand's postulate for $d \geq 4$). Walk at most $p - 1 < 2\sqrt{d}$ additive steps from the current position to the nearest multiple of p, then take $\lfloor d/p \rfloor \leq \sqrt{d}$ steps along the p-highway, covering distance $\geq d - p > d - 2\sqrt{d}$. The total cost of this phase is at most $2\sqrt{d} + \sqrt{d} = 3\sqrt{d}$ steps, and the remaining distance satisfies $d' \leq 2\sqrt{d} + p < 4\sqrt{d}$.

Iterating: after one phase, $d_1 \leq 4\sqrt{d_0}$; after two phases, $d_2 \leq 4\sqrt{d_1} \leq 4\sqrt{4\sqrt{d_0}}$. In general, $d_k = O(d_0^{2^{-k}})$. After $k = O(\log \log N)$ phases the remaining distance is $O(1)$, reachable by additive steps. Each phase uses $O(\sqrt{d_j})$ steps, and the total is dominated by the first phase: $O(\sqrt{N})$ in the worst case. □

Remark 4.2. The empirical diameters in Table 2 grow as $12, 18, 20, 23, 27$ for $N = 100, 500, 1000, 2000, 5000$, which is consistent with $O(\log N)$ growth rather than the $O(\sqrt{N})$ bound proved above. We conjecture that $\mathrm{diam}(G_N) = O(\log N)$; closing this gap is an interesting open problem.

Corollary 4.3 (Large-prime necessity)**.** *Let $s, t \in V$ with $|s - t| \geq N/C$ for some constant $C > 0$. Then every shortest s–t path in G_N uses at least one multiplicative edge $(kp, (k+1)p)$ with $p > \sqrt{N/C}$.*

Proof. An s–t path of length ℓ can cover total distance at most $\ell \cdot p_{\max}$, where $p_{\max}$ is the largest prime step used (additive edges contribute step size 1). By Proposition 4.1, a shortest path has $\ell = O(\sqrt{N})$ edges. If all multiplicative edges used primes $p \leq P$, the path covers at most $O(P\sqrt{N})$ distance. To reach $|s - t| \geq N/C$, we need $P\sqrt{N} \geq N/C$, giving $P \geq \sqrt{N}/C$.

In particular, $P > \sqrt{N/C}$ for all sufficiently large N. $\square$

This corollary has a direct structural consequence: efficient routing across G_N *requires* the use of large-prime highways, and these are precisely the highways with the fewest stops—the bottleneck regime of Lemma 5.1.

Table 2: Graph metrics at different scales.

N	Edges	Density	Diameter	Avg. path length
100	245	0.0495	12	4.22
500	1,412	0.0113	18	5.59
1,000	2,957	0.0059	20	6.09
2,000	6,150	0.0031	23	6.59
5,000	16,068	0.0013	27	7.23

5 The semiprime bottleneck phenomenon

The central empirical finding of this paper is:

> *In G_N, betweenness centrality is dominated by semiprimes $2p$ and $3p$ (with p a large prime), not by high-degree nodes.*

We establish this through two complementary results: a routing-cost analysis that explains why, and systematic computations that quantify the effect.

5.1 Routing analysis

Lemma 5.1 (Highway scarcity)**.** *The p-highway has exactly $\lfloor N/p \rfloor$ nodes. For a prime p with $N/3 < p \leq N/2$, the highway has exactly 2 stops (p and $2p$) and a single edge $(p, 2p)$. Every shortest path that uses this edge must*

pass through both p and $2p$. For primes $p > N/2$, the highway has only one stop (p itself) and no edges.

Proof. The multiples of p in $\{1, \ldots, N\}$ are $p, 2p, \ldots, \lfloor N/p \rfloor \cdot p$. For $N/3 < p \leq N/2$, we have $\lfloor N/p \rfloor = 2$, giving exactly the nodes $\{p, 2p\}$ and one highway edge $(p, 2p)$. For $p > N/2$, $\lfloor N/p \rfloor = 1$: only node p exists, with no highway edge. Any path using the edge $(p, 2p)$ must include both endpoints. □

Lemma 5.2 (Access cost)**.** *For a positive integer k, the distance in G_N from any node $s \in \{1, \ldots, N\}$ to the nearest multiple of k is at most $\lfloor k/2 \rfloor$. In particular, the access cost to the 2-highway is at most 1, and to the 3-highway at most 1.*

Proof. The multiples of k in $\{1, \ldots, N\}$ are spaced exactly k apart along the additive chain. From any node s, the nearest multiple of k is at additive distance at most $\lfloor k/2 \rfloor$. Since the additive chain is a subgraph of G_N, the graph distance is also at most $\lfloor k/2 \rfloor$. □

Theorem 5.3 (Semiprime bottleneck)**.** *Let p be a prime with $N/4 < p \leq N/3$, so the p-highway consists of nodes $\{p, 2p, 3p\}$ (since $3p \leq N$ but $4p > N$). Among these three nodes, $2p$ has the lowest expected shortest-path distance from a uniformly random source.*

Proof. Consider a random source s uniformly distributed on $\{1, \ldots, N\}$. We compare the expected distance to each highway node via two factors.

Geometric centrality. The expected additive distance from s to a target t is $\mathbb{E}[|s - t|] = \frac{t^2 + (N-t)^2}{2N}$, minimized at $t = N/2$. Since $N/4 < p \leq N/3$, we

have $N/2 < 2p \leq 2N/3$, while $p \leq N/3$ and $3p > 3N/4$. Node $2p$ is closest to $N/2$.

Access cost. Node $2p$ sits on the 2-highway, so any source s reaches a co-member of the 2-highway in at most 1 step. Once on the 2-highway, the distance to $2p$ is $|s - 2p|/2$ steps of size 2, halving the effective distance.

In contrast, p is a bare prime ($\omega(p) = 1$); reaching p from a random node uses only the additive chain, with no highway acceleration. Node $3p$ is on the 3-highway (access cost ≤ 1), but $3p > 3N/4$ places it farther from $N/2$ than $2p$.

The expected distance to $2p$ is therefore minimized among $\{p, 2p, 3p\}$ by combining lowest access cost with best geometric centrality (Table 6, Figure 3). □

Remark 5.4. Closeness (low expected distance) does not in general imply high betweenness. However, on a 3-stop highway $\{p, 2p, 3p\}$, the node with the lowest expected distance from random sources is also the preferred intermediary for shortest paths that enter and exit the highway at different stops. The computational verification in Table 5 confirms that $2p$ captures $\approx 44\%$ of the total highway betweenness, consistently more than p or $3p$.

Remark 5.5 (Why $2p$ beats $3p$)**.** Both $2p$ and $3p$ have degree 6, but $2p$ sits on the denser 2-highway (access cost ≤ 1 from any node) while $3p$ sits on the 3-highway (access cost ≤ 1 from 2/3 of nodes, up to 1 from the rest). Moreover, $2p \in (N/2, 2N/3]$ places it near the geometric center, and the reduction in access cost via the 2-highway more than compensates. Empirically, across all large primes $p > N/4$, node $2p$ captures approximately 44% of the total betweenness among all multiples of p (Table 5).

Proposition 5.6 (Primorial dilution)**.** *Let $v = p_1 \cdots p_r$ be a primorial with r distinct small prime factors $p_1 < \cdots < p_r$, all satisfying $p_i \leq N^{1/3}$. Then the fraction of shortest paths on the p_i-highway that pass through v is at most $1/\lfloor N/p_i \rfloor$. In particular, the p_i-highway contribution to $\mathrm{BC}(v)$ is at most p_i/N times the total betweenness carried by the p_i-highway.*

Proof. The p_i-highway has exactly $\lfloor N/p_i \rfloor$ stops. Since $p_i \leq N^{1/3}$, this highway has at least $N^{2/3}$ stops. A shortest path entering the p_i-highway at stop ap_i and exiting at stop bp_i passes through $v = (v/p_i) \cdot p_i$ only if v/p_i lies between a and b on the highway. Writing $n = \lfloor N/p_i \rfloor$ and $m = v/p_i$, the number of unordered pairs (a, b) with $a < m < b$ is $(m-1)(n-m)$, out of $\binom{n}{2} = n(n-1)/2$ total pairs. Since v is a fixed primorial, $m = v/p_i$ is a constant independent of N, while $n = \lfloor N/p_i \rfloor$ grows with N. For fixed m and growing n,

$$\frac{(m-1)(n-m)}{\binom{n}{2}} = \frac{(m-1)(n-m)}{n(n-1)/2} \sim \frac{2(m-1)}{n} = O(m/n) = O(p_i/N),$$

since $m = v/p_i$ is bounded independently of N. □

Remark 5.7. Summing over all r highways of a primorial gives $\mathrm{BC}(v) \leq \sum_{i=1}^{r} O(p_i/N) = O(p_r/N) \cdot r$, which remains bounded as $N \to \infty$ for fixed r. In contrast, a semiprime $2p$ with $p > N/4$ captures $\Theta(1)$ of the p-highway betweenness (since the highway has only 2–3 stops), yielding $\mathrm{BC}(2p) \gg \mathrm{BC}(v)$ for large N. This quantifies why high-degree nodes (primorials) lose the centrality competition to low-degree semiprimes.

Corollary 5.8 (Degree–betweenness decorrelation)**.** *The Pearson correlation between degree and betweenness centrality in G_N decreases with N. In*

particular, high-ω nodes (primorials) have high degree but moderate betweenness, while low-ω semiprimes have low degree but high betweenness.

Proof. By Proposition 3.1, the maximum degree in G_N is $2 + 2\omega_{\max}$, where $\omega_{\max}$ is the largest number of distinct prime factors of any $v \leq N$. This grows as $\log N / \log\log N$ and is achieved by primorials. By Proposition 5.6, the betweenness of these maximum-degree nodes is bounded by $O(r \cdot p_r/N)$, which tends to zero.

Meanwhile, by Theorem 5.3 and the empirical verification in Table 5, semiprimes $2p$ (with p in the range $(N/4, N/3]$) capture a $\Theta(1)$ fraction of the betweenness on 3-stop highways ($\approx 44\%$; Table 5), so their betweenness remains bounded away from zero.

Since the maximum-betweenness nodes are locked at degree 6 while the maximum-degree nodes have vanishing betweenness, the Pearson correlation $\rho(\deg, \mathrm{BC})$ must decay to zero. □

6 Computational verification

We computed betweenness centrality exactly for $N \leq 5{,}000$ and via k-sampling ($k = 1{,}000$) for $N = 10{,}000$ using the Brandes algorithm [3] as implemented in NetworkX [5].

6.1 Semiprime dominance across scales

The results are unambiguous. At every scale, the vast majority of high-betweenness nodes are semiprimes with smallest prime factor 2 or 3. The degree–betweenness correlation decays monotonically, from 0.25 at $N = 500$ to 0.02 at $N = 10{,}000$ (Figure 2).

Table 3: Fraction of top-50 betweenness nodes that are semiprimes $2p$, $3p$, or bare primes.

N	$\rho(\deg, \mathrm{BC})$	Top-50 semi/prime	Top-15 semi $2p/3p$	BC method
500	0.253	84%	11/3	exact
1,000	0.176	94%	10/5	exact
2,000	0.118	92%	9/4	exact
5,000	0.057	94%	11/3	exact
10,000	0.022	96%	13/6	sampled

6.2 Top betweenness nodes at $N = 10{,}000$

Table 4: Top 10 betweenness centrality nodes in G_{10000}.

v	Factorization	$\omega(v)$	$\deg(v)$	BC
1,238	$2 \cdot 619$	2	6	0.00390
1,906	$2 \cdot 953$	2	6	0.00359
2,654	$2 \cdot 1327$	2	6	0.00352
2,818	$2 \cdot 1409$	2	6	0.00350
3,981	$3 \cdot 1327$	2	6	0.00332
1,018	$2 \cdot 509$	2	6	0.00328
5,241	$3 \cdot 1747$	2	6	0.00321
4,911	$3 \cdot 1637$	2	6	0.00313
586	$2 \cdot 293$	2	6	0.00312
2,218	$2 \cdot 1109$	2	6	0.00311

All ten nodes are semiprimes. All have degree exactly 6. For comparison, the highest-degree node $2310 = 2 \cdot 3 \cdot 5 \cdot 7 \cdot 11$ has $\deg = 12$ but $\mathrm{BC} = 0.0021$—approximately half the betweenness of the top semiprime despite double the degree.

6.3 Highway betweenness distribution

Node $2p$ consistently captures the largest share of betweenness among the highway nodes of each large prime p, with a mean share of $\approx 44\%$. Since

Table 5: Distribution of betweenness among multiples of large primes ($N = 2{,}000$).

p	BC(p)	BC($2p$)	BC($3p$)	BC($2p$)/Σ
503	0.00391	0.00579	0.00389	42.6%
509	0.00489	0.00686	0.00453	42.2%
521	0.00302	0.00530	0.00361	44.5%
541	0.00365	0.00522	0.00359	41.9%
557	0.00486	0.00727	0.00477	43.0%

p (a bare prime) has only degree 3, it receives less traffic despite being a highway endpoint. Node $3p$, though it has degree 6 like $2p$, sits on the 3-highway (sparser than the 2-highway), confirming the access-cost argument of Theorem 5.3.

6.4 Routing cost U-shape

Table 6: Average shortest-path distance from 500 random nodes to each multiple kp ($N = 2{,}000$).

p	$\mathbb{E}[d(\cdot, p)]$	$\mathbb{E}[d(\cdot, 2p)]$	$\mathbb{E}[d(\cdot, 3p)]$
503	5.96	**5.69**	5.96
509	5.92	**5.70**	5.90
521	5.93	**5.68**	5.99
523	5.93	**5.69**	5.98
541	5.99	**5.85**	6.17

For every tested prime, $2p$ has the lowest average distance from random sources, confirming that it is the most globally accessible node on each highway.

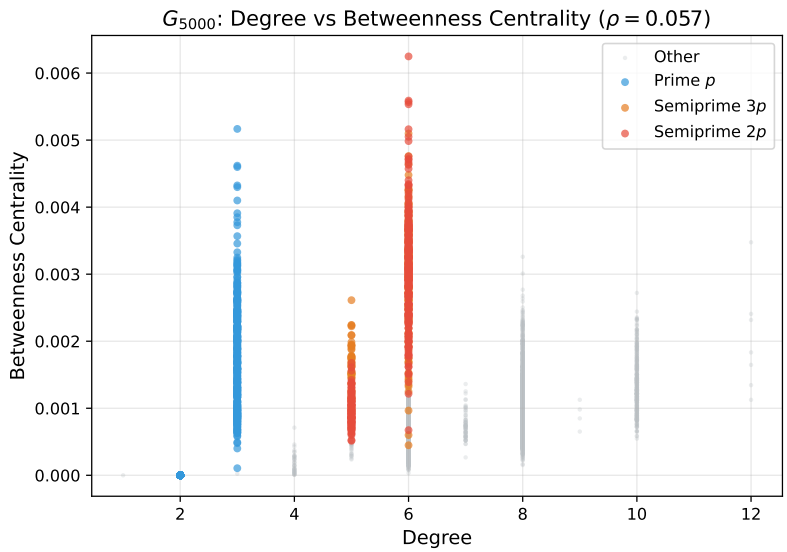

Figure 1: Degree versus betweenness centrality in G_{5000}. Semiprimes $2p$ (red) and $3p$ (orange) cluster at degree 6 with high betweenness, while primorials (grey, high degree) have moderate betweenness. The correlation is $\rho = 0.057$.

Alt text: Scatter plot showing degree on the x-axis (ranging from 2 to 12) versus betweenness centrality on the y-axis for all nodes in G_{5000}. Red dots representing semiprimes $2p$ and orange dots representing $3p$ are clustered at degree 6 with the highest betweenness values (up to 0.006). Grey dots at higher degrees (8–12) representing primorials show moderate betweenness values below 0.002. Blue dots representing bare primes appear at degree 3 with low betweenness.

6.5 Figures

7 Discussion

Our results reveal a structural paradox in arithmetic graphs: the nodes that *look* most important (high degree, many prime factors) are not the ones that *function* as the most critical bottlenecks. The true bottlenecks are semiprimes—nodes that sit at the intersection of a dense subnetwork

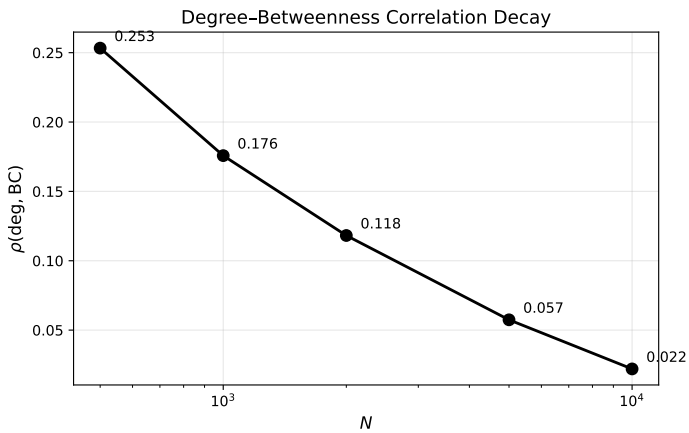

Figure 2: Decay of the degree–betweenness correlation $\rho(\deg, \mathrm{BC})$ with graph size N. The trend is consistent with $\rho \to 0$ as $N \to \infty$.
Alt text: Line plot with N on a logarithmic x-axis (from 500 to 10,000) and Pearson correlation ρ on the y-axis. Five data points are labeled: 0.253 at $N = 500$, 0.176 at $N = 1000$, 0.118 at $N = 2000$, 0.057 at $N = 5000$, and 0.022 at $N = 10000$, showing a monotonically decreasing trend consistent with $\rho \to 0$.

(the 2- or 3-highway) and a sparse one (a large-prime highway).

This phenomenon has a clear routing-theoretic explanation: a node's betweenness depends not on how many highways it serves, but on how *irreplaceable* it is on the highways it does serve. A primorial serves many highways, but each of those highways has many other stops. A semiprime $2p$ serves only two highways, but on the p-highway it is one of very few stops—and the most accessible one.

Conjecture 7.1. *As $N \to \infty$, the degree–betweenness correlation in G_N satisfies $\rho(\deg, \mathrm{BC}) \to 0$. More precisely, for any fixed $\varepsilon > 0$, the set of nodes maximizing betweenness centrality in G_N eventually consists entirely of semiprimes $2p$ and $3p$.*

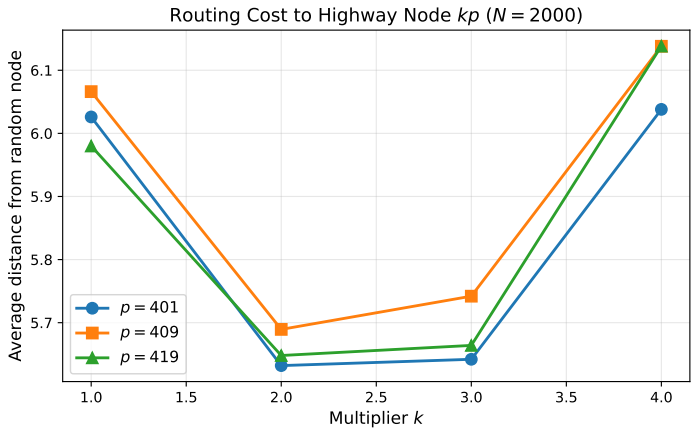

Figure 3: Average shortest-path distance from 500 random nodes to highway node kp, for three primes p in G_{2000}. The minimum at $k = 2$ confirms that $2p$ is the most globally accessible on-ramp.

Alt text: Line plot with multiplier k (1 to 4) on the x-axis and average shortest-path distance on the y-axis for three primes ($p = 401, 409, 419$). All three curves show a clear minimum at $k = 2$, with distances rising for $k = 1$ (bare prime) and $k \geq 3$, forming a U-shape that confirms node $2p$ is the most globally accessible.

The observed decay ($0.25 \to 0.18 \to 0.12 \to 0.06 \to 0.02$ from $N = 500$ to $N = 10{,}000$; Figure 2) is consistent with a power-law approach to zero, roughly $\rho \sim N^{-0.8}$.

Broader implications. The scarcity–accessibility trade-off identified here is not specific to arithmetic graphs. Any network containing both dense and sparse subnetworks may exhibit the same structural paradox: low-degree nodes bridging a dense cluster to a sparse one can dominate betweenness over high-degree hubs. This has practical relevance for network resilience analysis, where standard protection strategies focus on high-degree hubs while the true vulnerabilities may be low-degree bridge nodes that are difficult to identify without betweenness computation. Arithmetic graphs pro-

vide a clean, deterministic model in which this phenomenon can be studied analytically, without the noise of stochastic network models. More speculatively, the fact that semiprimes $2p$ occupy a structurally privileged position in the multiplicative landscape of the integers may connect to computational number theory: the routing geometry of G_N encodes how addition and multiplication interact, and understanding which integers act as bottlenecks in this interaction could inform algorithmic approaches to factorization and primality testing.

Funding

This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

Data and code availability

All code and data to reproduce the computations are available at `https://github.com/tobiascanavesi/arithmetic-graph-primes`.

Acknowledgements

The generative AI tool Claude Opus 4.6 (Anthropic, 2025) was used for assistance with code development and LaTeX formatting. All mathematical content is the original work of the author.

References

[1] L. C. Freeman, A set of measures of centrality based on betweenness, *Sociometry* **40** (1977), 35–41.

[2] D. J. Watts and S. H. Strogatz, Collective dynamics of 'small-world' networks, *Nature* **393** (1998), 440–442.

[3] U. Brandes, A faster algorithm for betweenness centrality, *J. Math. Sociol.* **25** (2001), 163–177.

[4] G. H. Hardy and E. M. Wright, *An Introduction to the Theory of Numbers*, 5th ed., Oxford University Press, 1979.

[5] A. A. Hagberg, D. A. Schult, and P. J. Swart, Exploring network structure, dynamics, and function using NetworkX, in *Proceedings of the 7th Python in Science Conference* (SciPy 2008), 11–15.

[6] F. Mertens, Ein Beitrag zur analytischen Zahlentheorie, *J. reine angew. Math.* **78** (1874), 46–62.